

# UofM Fintech Group Project 2

LSTM, GRU with Ethereum price prediction and forecast.

we will see if you gonna present

or how you going to present lol

kyle plathe

I'll do my best



KP

by Kyle Plathe, Richard Melvin & Meek Msaki, 2022

# Table of Contents

- Hypothesis
- Data Collection - rich, meek, kyle
- Data Cleanup - rich, meek, kyle
- LSTM model - meek
- GRU Model - kyle
- PCA - rich
- Discussion & Summary
- Postmortem
- Questions

# Hypothesis

# Hypothesis

*Our motivation and summary...*

- We use machine learning tools to predict the price of ethereum from historical data, economic indicators, and community sentiment on ethereum specifically from twitter.
- We test this hypothesis by building LSTM and GRU models

# 1. Data Collection

Describing what kinds of data we needed and where we find it.

# Our Sources

- Rich - Twint Protocol - Collecting tweet data
- Meek - Owlracle API - Getting Gas price history
- Kyle - Kaggle - ETH to USD Historical Data
- Kyle - FRED - Personal Savings Percentage
- Kyle - Market Watch - S&P 500 Historical Data







## 2. Data Cleanup

Preparing our data for our models

# Twint Protocol

Richard Melvin

- **Purpose:** Obtain data for testing the hypothesis that community sentiment predicts Ethereum price.
- **Method:** "Scrape" Twitter for tweets containing keyword, "Ethereum".
  - Twint project: <https://github.com/twintproject/twint>
- **Result:** 4.3 million tweets were collected (Nov 2017 to Dec 2020)







# Twint Protocol Flow Chart

**Twint**  
Install from github  
Requires: nest\_asyncio

```
git clone --depth=1
https://github.com/twintproject/twint.git
cd twint
pip install . -r requirements.txt
pip install nest_asyncio
```

**Packages**  
twint  
nest\_asyncio  
pandas

```
import twint
import nest_asyncio
nest_asyncio.apply()
import pandas as pd
```

**Configure Twint**

```
c = twint.Config()
```

**Search**  
"Ethereum"  
Nov '17 to Dec '20  
English language  
Pandas format

```
c.Search = "Ethereum"
c.Since = '2017-11-01'
c.Until = '2020-12-31'
c.Lang = 'en'
c.Pandas = TRUE
```

**Dataframe**  
date  
tweet

```
def twint_to_pd(columns):
    return twint.output.pandas.Tweets_df[columns]

twint_to_pd(["date", "tweet"])
```



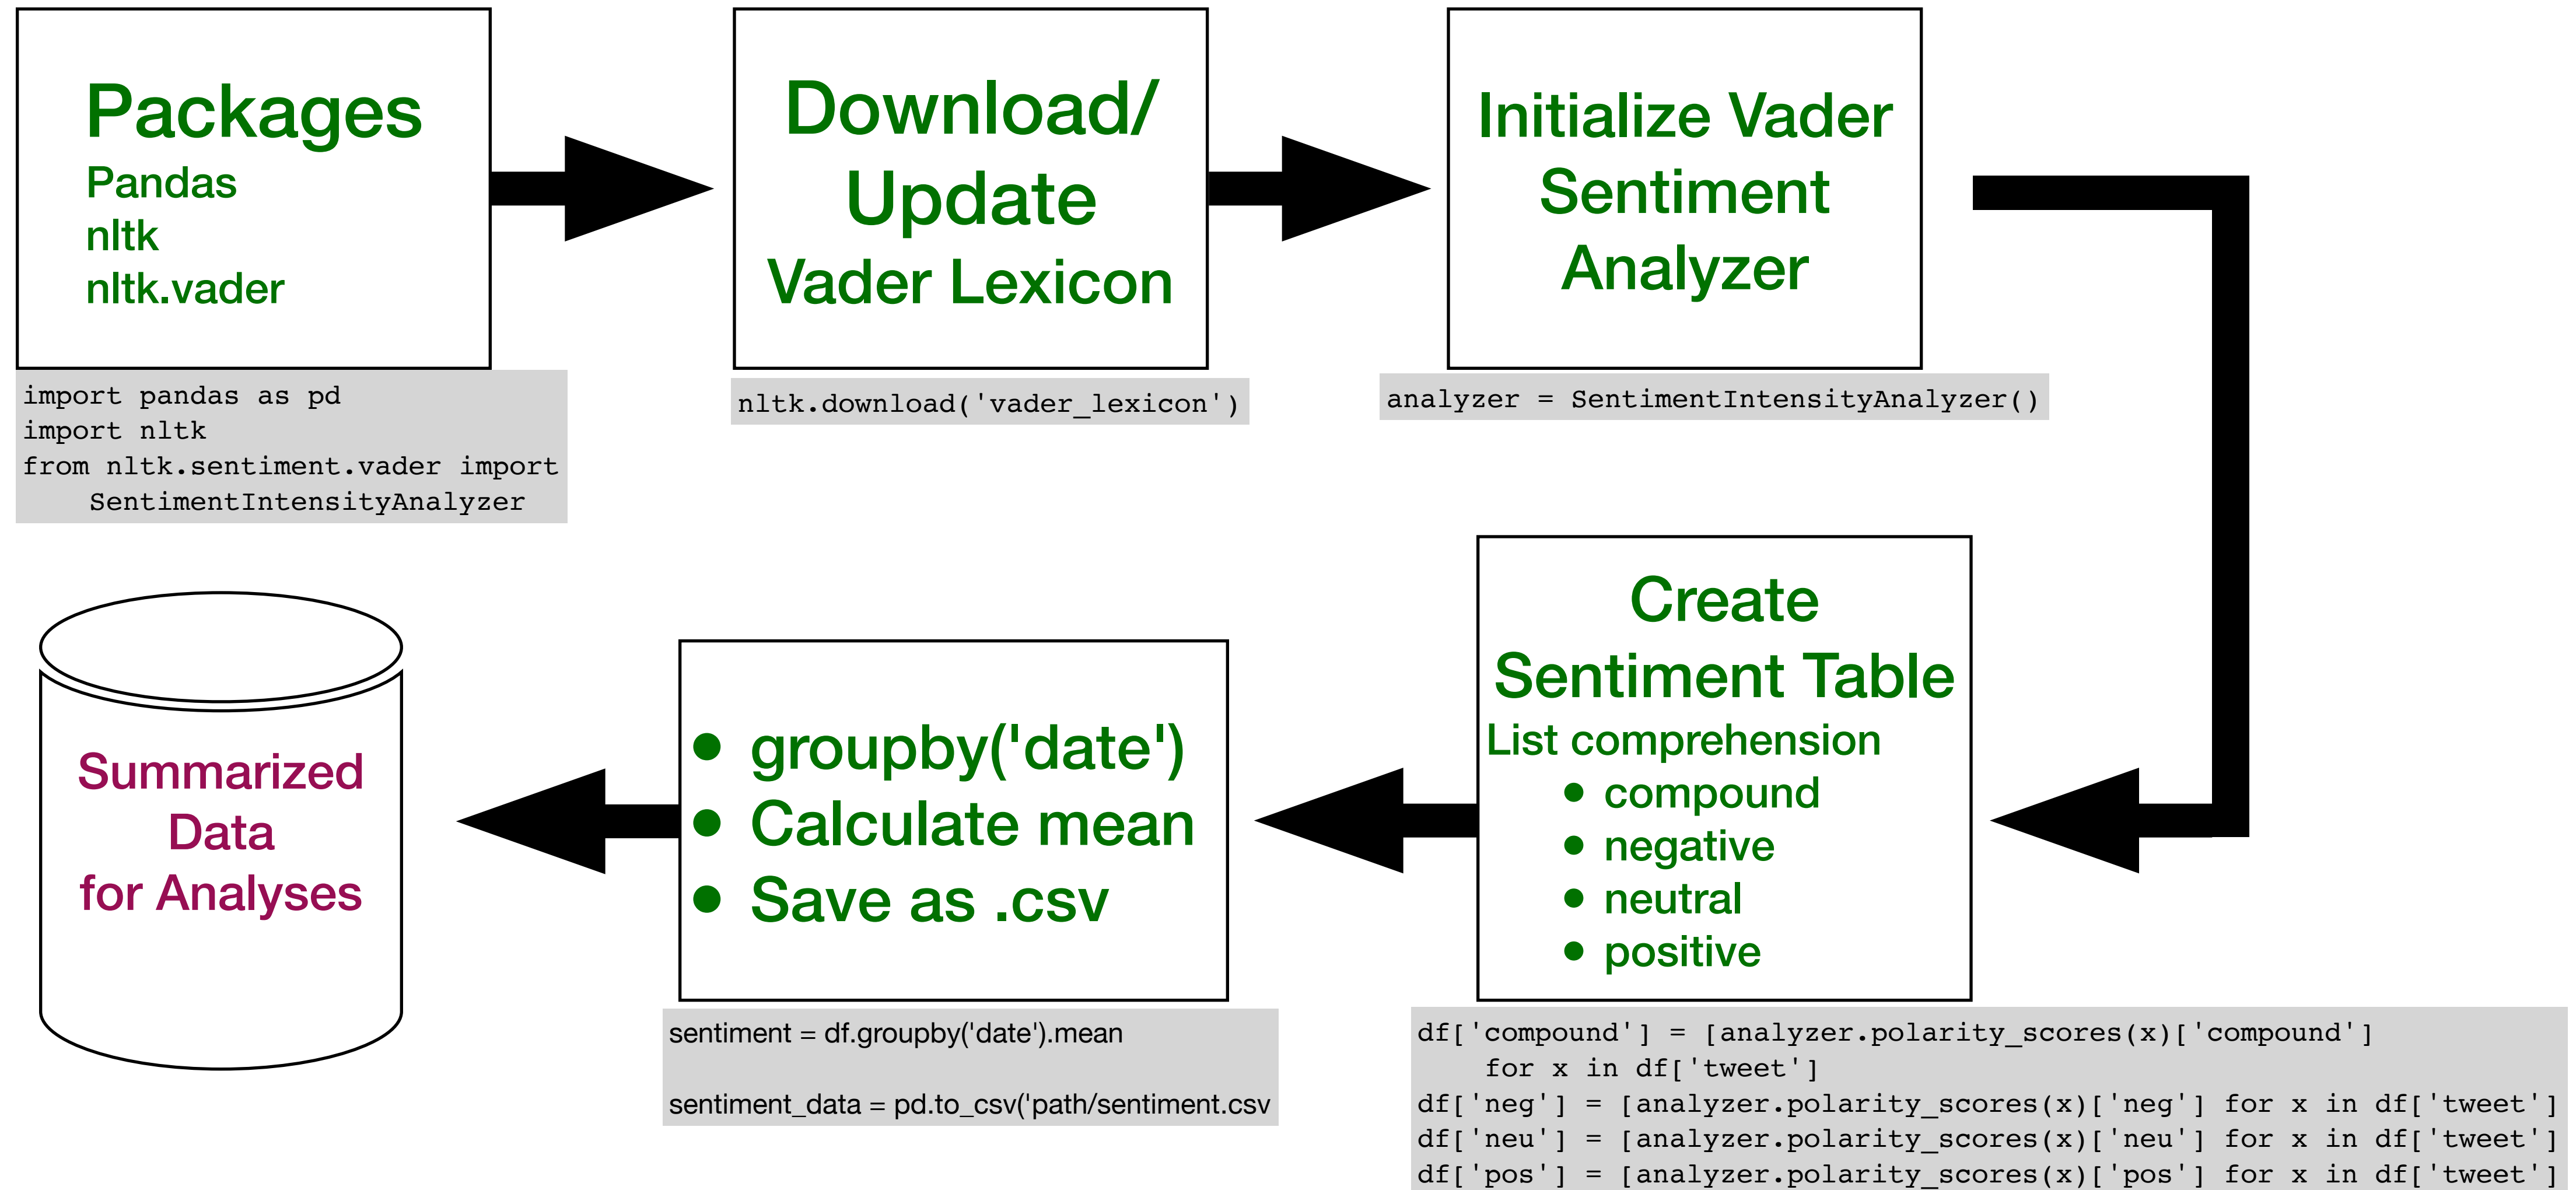
- Time consuming.
  - connection breaks.
  - restarts.
- Unlimited data
- Iterative function will improve.

# 3. Sentiment Analysis

- **Purpose:** Extract sentiment from Tweet texts.
- **Method:** Natural Language Toolkit (*nltk*), ***Vader*** sentiment analysis
- **Result:** Sentiment of 4.3 million tweets was quantified and summarized as the mean for each day.
  - Degree of positive or negative sentiment.



# Vader Protocol Flowchart



# Kyle's Data

- Using CVS files data from Kaggle ETH historical data, data from FRED - Personal Savings Percentage and historical data from the S&P 500 from Market Watch. I concatenated the 3 sets of data.
- One problem was that the Savings Percentage data was missing days from the month.

Using resample “D” for day and ffill for forward fill.

```
# Fill in the missing data days of the month  
Personal_Daily_Savings = Personal_Savings_df.resample("D").ffill()  
Personal_Daily_Savings.head()
```

Personal Saving %	
Date	
2014-07-01	7.2
2014-07-02	7.2
2014-07-03	7.2
2014-07-04	7.2
2014-07-05	7.2

# 4. LSTM Model - Meek

Using our data with LSTM (Long Short-Term Memory) model



```
[1]: # Initial imports
import pandas as pd
import numpy as np
import hvplot.pandas

# Import dependencies
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense, Dropout
from sklearn.preprocessing import MinMaxScaler
from tensorflow.keras import layers
```

our imports

2

gas price historical  
data from csv

```
# Set the random seed for reproducibility
# Note: This is for the homework solution, but it is good practice to comment_
↳ this out and run multiple experiments to evaluate your model
from numpy.random import seed
seed(72)
from tensorflow import random
random.set_seed(72)
```

```
[ ]: # Upload eth_gas_prices.csv, eth_features, eth_sentiment_history data to Colab
from google.colab import files
csv_file = files.upload()
```

<IPython.core.display.HTML object>

```
[2]: # Read in data
gas_df = pd.read_csv("eth_gas_prices_history.csv", index_col = "date",
↳ infer_datetime_format=True, parse_dates=True)

# Set index to a datetime format
gas_df.index = pd.to_datetime(gas_df.index)

# # View first 5 rows of gas_df
gas_df.head()
```

```
[2]:
```

	open	close	low	high	avgGas
date					
2022-07-05	22.37	30.70	10.92	92.03	89140.499588
2022-07-04	30.92	21.04	8.46	159.32	89253.609177
2022-07-03	14.06	42.13	7.20	76.55	92493.795505
2022-07-02	50.78	32.84	5.00	320.81	92896.649370
2022-07-01	12.74	47.59	7.49	103.23	91907.403785

Eth price historical  
+  
other features  
data from csv

```
[3]: # Load the historical closing prices for Bitcoin
eth_df = pd.read_csv('eth_features.csv', index_col="Date",
↳ infer_datetime_format=True, parse_dates=True)

# Sort index
eth_df = eth_df.sort_index(ascending=False)
eth_df = pd.DataFrame(eth_df, index = eth_df.index)
eth_df.index = eth_df.index.strftime("%Y-%m-%d")

# Set index to a datetime format
eth_df.index = pd.to_datetime(eth_df.index)

eth_df.head()
```



[3]:

	Open	High	Low	Close	Volume \
Date					
2022-03-25	3109.523438	3182.826660	3097.624268	3122.535889	16882068480
2022-03-24	3031.060791	3118.387695	3012.326660	3108.062012	18070503166
2022-03-23	2973.145020	3036.752197	2933.306641	3031.067139	16008767658
2022-03-22	2897.774170	3040.382813	2892.544434	2973.131104	16830539230
2022-03-21	2860.103271	2954.556641	2838.250488	2897.976563	15206116098

	Personal Saving %	Close/Last
Date		
2022-03-25	5.0	4543.06
2022-03-24	5.0	4520.16
2022-03-23	5.0	4456.24
2022-03-22	5.0	4511.61
2022-03-21	5.0	4461.18

```
[4]: # Read eth_sentiment_history
sentiment_df = pd.read_csv('eth_sentiment_history.csv', index_col="date",
    infer_datetime_format=True, parse_dates=True)
sentiment_df = sentiment_df.drop(columns=["Day of Year", "id", "Unnamed: 0"])

sentiment_df
```

[4]:

	nlikes	nretweets	compound	neg	neu	pos
date						
2020-01-01	3.508475	0.855932	0.097053	0.032445	0.905449	0.062110
2020-01-02	4.185213	0.843540	0.101582	0.052678	0.872564	0.074755
2020-01-03	2.815929	0.701327	0.105520	0.039138	0.892830	0.068028
2020-01-04	2.978692	0.728444	0.086918	0.049989	0.880994	0.069022
2020-01-05	2.992630	1.079687	0.104279	0.045839	0.886108	0.068048
...	...	...	...	...	...	...
2017-12-25	2.453402	1.475008	0.193739	0.021307	0.885513	0.093183
2017-12-26	1.373333	1.226667	0.152753	0.015047	0.915843	0.069110
2017-12-27	3.701015	1.931156	0.149329	0.020644	0.908841	0.070512
2017-12-29	2.524540	0.911043	0.222387	0.018279	0.887650	0.094071
2017-12-30	1.250290	0.890309	0.188613	0.019919	0.896495	0.083584

[1049 rows x 6 columns]

```
[5]: # Join historical data for ethereum, gas prices and twitter sentiment with
    highest gas fees
df = eth_df.join([gas_df, sentiment_df], how="inner")

# Sort index
df = df.sort_index(ascending=False)
df
```

[5]:

	Open	High	Low	Close	Volume \
2020-12-30	731.472839	754.303223	720.988892	751.618958	17294574210
2020-12-29	730.358704	737.952881	692.149414	731.520142	18710683199
2020-12-28	683.205811	745.877747	683.205811	730.397339	24222565862
2020-12-24	584.135620	613.815186	568.596375	611.607178	14317413703
2020-12-23	634.824585	637.122803	560.364258	583.714600	15261413038
...	...	...	...	...	...
2017-11-15	337.963989	340.911987	329.812988	333.356995	722665984
2017-11-14	316.763000	340.177002	316.763000	337.631012	1069680000
2017-11-13	307.024994	328.415009	307.024994	316.716003	1041889984
2017-11-10	320.670990	324.717987	294.541992	299.252991	885985984
2017-11-09	308.644989	329.451996	307.056000	320.884003	893249984

	Personal Saving %	Close/Last	open	close	low	high \
2020-12-30	14.0	3732.04	47.15	74.0	1.0	205.0
2020-12-29	14.0	3727.04	68.31	68.2	9.0	178.0
2020-12-28	14.0	3735.36	55.00	45.0	9.0	153.0
2020-12-24	14.0	3703.06	40.00	34.7	1.0	202.0
2020-12-23	14.0	3690.01	33.00	88.0	1.0	579.0
...	...	...	...	...	...	...
2017-11-15	7.0	2564.62	0.10	20.0	0.1	60.0
2017-11-14	7.0	2578.87	30.00	1.0	0.1	35.0
2017-11-13	7.0	2584.84	4.00	1.0	0.1	50.0
2017-11-10	7.0	2582.30	8.00	0.0	0.0	34.0
2017-11-09	7.0	2584.62	4.00	20.0	0.1	60.0

	avgGas	nlikes	nretweets	compound	neg	neu \
2020-12-30	76337.160077	5.751632	1.361131	0.215897	0.023943	0.880747
2020-12-29	72782.380459	5.482190	1.015992	0.181098	0.029384	0.883803
2020-12-28	69025.670836	6.617657	1.088464	0.159243	0.035662	0.876452
2020-12-24	76968.816529	5.794193	1.053289	0.197436	0.030146	0.875286
2020-12-23	71480.326482	6.123096	1.344860	0.203225	0.032519	0.871497
...	...	...	...	...	...	...
2017-11-15	57898.769541	3.261106	2.833403	0.098748	0.029318	0.903320
2017-11-14	77198.609591	2.898046	2.153568	0.155619	0.030826	0.877397
2017-11-13	68732.650465	3.146906	3.628248	0.103907	0.029216	0.901205
2017-11-10	62459.741726	2.460854	2.444617	0.074516	0.053909	0.869239
2017-11-09	51382.905682	2.580446	2.625297	0.086666	0.041577	0.887715

	pos
2020-12-30	0.095315
2020-12-29	0.086815
2020-12-28	0.087886
2020-12-24	0.094560
2020-12-23	0.095986
...	...
2017-11-15	0.067360

Joining all dataframes

- eth\_df
- gas\_df
- sentiment\_df

```
2017-11-14  0.091771
2017-11-13  0.069577
2017-11-10  0.076842
2017-11-09  0.070710
```

```
[691 rows x 18 columns]
```

```
[6]: # This function accepts the column number for the features (X) and the target
      ↪ (y)
      # It chunks the data up with a rolling window of Xt-n to predict Xt
      # It returns a numpy array of X any y
      def window_data(df, window, feature_col_number, target_col_number):
          X = []
          y = []
          for i in range(len(df) - window - 1):
              features = df.iloc[i:(i + window), feature_col_number]
              target = df.iloc[(i + window), target_col_number]
              X.append(features)
              y.append(target)
          return np.array(X), np.array(y).reshape(-1, 1)
```

```
[7]: # Predict Closing Prices using a 10 day window of previous closing prices
      # Then, experiment with window sizes anywhere from 1 to 10 and see how the
      ↪ model performance changes
      window_size = 10

      # Column index 0 is the 'Open' column for eth_df
      # Column index 3 is the 'close' column for eth_df
      feature_column = 3
      target_column = 3
      X, y = window_data(df, window_size, feature_column, target_column)
```

```
[8]: # Use 70% of the data for training and the remainder for testing
      split = int(0.7 * len(X))
      X_train = X[: split]
      y_train = y[: split]

      X_test = X[split:]
      y_test = y[split:]
```

```
[9]: from sklearn.preprocessing import MinMaxScaler
      # Use the MinMaxScaler to scale data between 0 and 1.
      scaler = MinMaxScaler()
      scaler.fit(X_train)

      # Scale thr X_train and X_test sets
      X_train = scaler.transform(X_train)
```

```
X_test = scaler.transform(X_test)

# fit the MinMaxScaler object with the target daya
scaler.fit(y_train)

# Scale the y_train and y_test sets
y_train = scaler.transform(y_train)
y_test = scaler.transform(y_test)
```

```
[10]: # Reshape the features for the model
      X_train = X_train.reshape((X_train.shape[0], X_train.shape[1], 1))
      X_test = X_test.reshape((X_test.shape[0], X_test.shape[1], 1))
```

```
[11]: # Build the LSTM model.
      # The return sequences need to be set to True if you are adding additional LSTM
      ↪ layers, but
      # You don't have to do this for the final layer.
      # Note: The dropouts help prevent overfitting
      # Note: The input shape is the number of time steps and the number of indicators
      # Note: Batching inputs has a different input shape of Samples/TimeSteps/
      ↪ Features
```

```
# Define model
lstm_model = Sequential()

# Initial model setup
number_units = 30
dropout_fraction = 0.2
```

```
# Layer 1
lstm_model.add(LSTM(
    units=number_units,
    return_sequences=True,
    input_shape=(X_train.shape[1], 1))
)
lstm_model.add(Dropout(dropout_fraction))
```

```
# Layer 2
lstm_model.add(LSTM(units=number_units, return_sequences=True))
lstm_model.add(Dropout(dropout_fraction))
```

```
# Layer 3
lstm_model.add(LSTM(units=number_units))
lstm_model.add(Dropout(dropout_fraction))
```

```
# Output layer
lstm_model.add(Dense(1))
```

selecting column 3  
as our column  
for prediction eth  
closing prices

LSTM  
layers

```
[12]: # Compile the model
lstm_model.compile(optimizer="adam", loss="mean_squared_error")
```

```
[13]: # Summarize the model
lstm_model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
lstm (LSTM)	(None, 10, 30)	3840
dropout (Dropout)	(None, 10, 30)	0
lstm_1 (LSTM)	(None, 10, 30)	7320
dropout_1 (Dropout)	(None, 10, 30)	0
lstm_2 (LSTM)	(None, 30)	7320
dropout_2 (Dropout)	(None, 30)	0
dense (Dense)	(None, 1)	31

=====  
Total params: 18,511  
Trainable params: 18,511  
Non-trainable params: 0  
=====

```
[14]: # Train the model
# Experiment with the batch size, but a smaller batch size is recommended
lstm_model.fit(X_train, y_train, epochs=10, shuffle=False, batch_size=100,
               verbose=1)
```

```
Epoch 1/10
5/5 [=====] - 4s 14ms/step - loss: 0.0912
Epoch 2/10
5/5 [=====] - 0s 14ms/step - loss: 0.0596
Epoch 3/10
5/5 [=====] - 0s 15ms/step - loss: 0.0368
Epoch 4/10
5/5 [=====] - 0s 17ms/step - loss: 0.0233
Epoch 5/10
5/5 [=====] - 0s 14ms/step - loss: 0.0208
Epoch 6/10
5/5 [=====] - 0s 15ms/step - loss: 0.0164
Epoch 7/10
```

**LSTM  
model  
evaluation  
91% accuracy**



```
5/5 [=====] - 0s 23ms/step - loss: 0.0128
Epoch 8/10
5/5 [=====] - 0s 27ms/step - loss: 0.0116
Epoch 9/10
5/5 [=====] - 0s 36ms/step - loss: 0.0097
Epoch 10/10
5/5 [=====] - 0s 26ms/step - loss: 0.0091
```

```
[14]: <keras.callbacks.History at 0x7fb009447250>
```

```
[15]: # Evaluate the model
lstm_model.evaluate(X_test, y_test, verbose = 0)
```

```
[15]: 0.08968404680490494
```

```
[16]: # Make some predictions
predicted = lstm_model.predict(X_test)
```

```
[17]: # Recover the original prices instead of the scaled version
predicted_prices = scaler.inverse_transform(predicted)
real_prices = scaler.inverse_transform(y_test.reshape(-1, 1))
```

```
[18]: # Create a DataFrame of Real and Predicted values
eth_price_prediction_lstm_df = pd.DataFrame({
    "real eth price": real_prices.ravel(),
    "predicted eth price": predicted_prices.ravel()
}, index = df.index[-len(real_prices): ])
eth_price_prediction_lstm_df.head()
```

```
[18]:
```


	real eth price	predicted eth price
2018-10-08	227.981995	208.017929
2018-10-05	229.255005	208.443634
2018-10-04	227.600998	209.198181
2018-10-03	222.218002	210.331802
2018-10-02	220.488998	211.785172

```
[19]: # Plot the real vs predicted eth values as a line chart
hvplot.extension('bokeh')

eth_price_prediction_lstm_df.hvplot.line(title = "actual eth price vs.
predicted eth prices with lstm").opts(width = 800)
```

```
[19]: :NdOverlay      [Variable]
      :Curve       [index]   (value)
```

**actual eth price vs. predicted eth prices with lstm**



# **5. GRU Model - KYLE**

Using our data with GRU (Gated Recurrent Unit) model



# ETH GRU Prediction

July 7, 2022

```
[1]: # Import dependencies
import numpy as np
from numpy import newaxis
import pandas as pd
from keras.layers.core import Dense, Activation, Dropout
from keras.layers.recurrent import LSTM, GRU
from keras.models import Sequential
from keras import optimizers
from sklearn.preprocessing import MinMaxScaler
import matplotlib.pyplot as plt
plt.style.use('fivethirtyeight')
```

```
[2]: # Enter in how many steps we will enroll the network.

Enrol_window = 100

print ('enroll window set to', Enrol_window )
```

enroll window set to 100

```
[3]: # Support functions
sc = MinMaxScaler(feature_range=(0,1))
def load_data(datasetname, column, seq_len, normalise_window):
    # A support function to help prepare datasets for an RNN/LSTM/GRU
    data = datasetname.loc[:,column]

    sequence_length = seq_len + 1
    result = []
    for index in range(len(data) - sequence_length):
        result.append(data[index: index + sequence_length])

    if normalise_window:
        #result = sc.fit_transform(result)
        result = normalise_windows(result)

    result = np.array(result)

    #Last 10% is used for validation test, first 90% for training
```

## Support functions:

- load\_data ( )
- normalise\_windows ( )
- predict\_sequence\_full ( )
- predict\_sequences\_multiple ( )
- plot\_results ( )
- plot\_results\_multiple ( )

```
row = round(0.9 * result.shape[0])
train = result[:int(row), :]
np.random.shuffle(train)
x_train = train[:, :-1]
y_train = train[:, -1]
x_test = result[int(row):, :-1]
y_test = result[int(row):, -1]

x_train = np.reshape(x_train, (x_train.shape[0], x_train.shape[1], 1))
x_test = np.reshape(x_test, (x_test.shape[0], x_test.shape[1], 1))

return [x_train, y_train, x_test, y_test]

def normalise_windows(window_data):
    # A support function to normalize a dataset
    normalised_data = []
    for window in window_data:
        normalised_window = [((float(p) / float(window[0])) - 1) for p in
window]
        normalised_data.append(normalised_window)
    return normalised_data

def predict_sequence_full(model, data, window_size):
    #Shift the window by 1 new prediction each time, re-run predictions on new
window
    curr_frame = data[0]
    predicted = []
    for i in range(len(data)):
        predicted.append(model.predict(curr_frame[newaxis,:,:])[0,0])
        curr_frame = curr_frame[1:]
        curr_frame = np.insert(curr_frame, [window_size-1], predicted[-1],
axis=0)
    return predicted

def predict_sequences_multiple(model, data, window_size, prediction_len):
    #Predict sequence of <prediction_len> steps before shifting prediction run
forward by <prediction_len> steps
    prediction_seqs = []
    for i in range(int(len(data)/prediction_len)):
        curr_frame = data[i*prediction_len]
        predicted = []
        for j in range(prediction_len):
            predicted.append(model.predict(curr_frame[newaxis,:,:])[0,0])
            curr_frame = curr_frame[1:]
            curr_frame = np.insert(curr_frame, [window_size-1], predicted[-1],
axis=0)
        prediction_seqs.append(predicted)
```



```

    return prediction_seqs

def plot_results(predicted_data, true_data):
    fig = plt.figure(facecolor='white')
    ax = fig.add_subplot(111)
    ax.plot(true_data, label='True Data')
    plt.plot(predicted_data, label='Prediction')
    plt.legend()
    plt.show()

def plot_results_multiple(predicted_data, true_data, prediction_len):
    fig = plt.figure(facecolor='white')
    ax = fig.add_subplot(111)
    ax.plot(true_data, label='True Data')
    #Pad the list of predictions to shift it in the graph to it's correct start
    for i, data in enumerate(predicted_data):
        padding = [None for p in range(i * prediction_len)]
        plt.plot(padding + data, label='Prediction')
        plt.legend()
    plt.show()

print ('Support functions defined')
```

Support functions defined

```
[15]: # Upload CSV file
from google.colab import files
uploaded = files.upload()
```

<IPython.core.display.HTML object>

Saving eth\_all\_features.csv to eth\_all\_features (3).csv

```
[26]: # Read Eth data
dataset = pd.read_csv('eth_all_features.csv', index_col='Unnamed: 0',
    parse_dates=True)
dataset.tail()
```

[26]:

	Open	High	Low	Close	Volume	\
2017-11-15	337.963989	340.911987	329.812988	333.356995	722665984	
2017-11-14	316.763000	340.177002	316.763000	337.631012	1069680000	
2017-11-13	307.024994	328.415009	307.024994	316.716003	1041889984	
2017-11-10	320.670990	324.717987	294.541992	299.252991	885985984	
2017-11-09	308.644989	329.451996	307.056000	320.884003	893249984	

	Personal Saving %	Close/Last	open	close	low	high	\
2017-11-15	7.0	2564.62	0.1	20.0	0.1	60.0	
2017-11-14	7.0	2578.87	30.0	1.0	0.1	35.0	

2017-11-13	7.0	2584.84	4.0	1.0	0.1	50.0	
2017-11-10	7.0	2582.30	8.0	0.0	0.0	34.0	
2017-11-09	7.0	2584.62	4.0	20.0	0.1	60.0	

	avgGas	nlikes	nretweets	compound	neg	neu	\
2017-11-15	57898.769541	3.261106	2.833403	0.098748	0.029318	0.903320	
2017-11-14	77198.609591	2.898046	2.153568	0.155619	0.030826	0.877397	
2017-11-13	68732.650465	3.146906	3.628248	0.103907	0.029216	0.901205	
2017-11-10	62459.741726	2.460854	2.444617	0.074516	0.053909	0.869239	
2017-11-09	51382.905682	2.580446	2.625297	0.086666	0.041577	0.887715	

	pos
2017-11-15	0.067360
2017-11-14	0.091771
2017-11-13	0.069577
2017-11-10	0.076842
2017-11-09	0.070710


```
[31]: # Prepare the dataset, note that the eth price data will be normalized between
    0 and 1
    # A label is the thing we're predicting
    # A feature is an input variable, in this case ethereum price
    # Selected 'Close' (eth price at closing) attribute for prices. Let's see what
    it looks like

feature_train, label_train, feature_test, label_test = load_data(dataset,
    'Close', Enrol_window, True)

dataset["Close"]['2019'].plot(figsize=(16,4),legend=True)
dataset["Close"]['2020:'].plot(figsize=(16,4),legend=True) # 20% is used for
    training data which is approx 2022 data
plt.legend(['Training set (First 80%, approx before 2021)', 'Test set (Last 20%,
    approx 2022 and beyond)'])
plt.title('Ethereum price')
plt.show()
```

/usr/local/lib/python3.7/dist-packages/ipykernel\_launcher.py:9: FutureWarning: Value based partial slicing on non-monotonic DatetimeIndexes with non-existing keys is deprecated and will raise a KeyError in a future Version.

```
if __name__ == '__main__':
```



```
[32]: # The same LSTM model I would like to test, lets see if the sinus prediction
      ↪ results can be matched

model = Sequential()
model.add(GRU(50, return_sequences=True, input_shape=(feature_train.
      ↪ shape[1],1)))
model.add(Dropout(0.2))
model.add(GRU(100, return_sequences=False))
model.add(Dropout(0.2))
model.add(Dense(1, activation = "linear"))

model.compile(loss='mse', optimizer='adam')

print ('model compiled')

print (model.summary())
```

```
model compiled
Model: "sequential_2"

-----
Layer (type)                Output Shape          Param #
-----
gru_3 (GRU)                  (None, 100, 50)      7800
dropout_4 (Dropout)          (None, 100, 50)      0
gru_4 (GRU)                  (None, 100)           45300
dropout_5 (Dropout)          (None, 100)           0
dense_2 (Dense)              (None, 1)             101
=====

Total params: 53,201
Trainable params: 53,201
Non-trainable params: 0
```


```
-----
None
```

```
[33]: #Train the model
      model.fit(feature_train, label_train, batch_size=512, epochs=5, validation_data=
      ↪ (feature_test, label_test))
```


```
Epoch 1/5
2/2 [=====] - 8s 937ms/step - loss: 1.0042 - val_loss: 0.3466
Epoch 2/5
2/2 [=====] - 1s 320ms/step - loss: 0.6146 - val_loss: 0.1761
Epoch 3/5
2/2 [=====] - 1s 309ms/step - loss: 0.3265 - val_loss: 0.0893
Epoch 4/5
2/2 [=====] - 1s 314ms/step - loss: 0.1323 - val_loss: 0.1065
Epoch 5/5
2/2 [=====] - 1s 290ms/step - loss: 0.0755 - val_loss: 0.2019
```

```
[33]: <keras.callbacks.History at 0x7fd6d570e290>
```

```
[34]: # Model and predict the eth price
      predicted_eth_price = model.predict(feature_test)
      plot_results(predicted_eth_price,label_test)
```



**actual eth price vs. predicted eth prices with GRU model**



```
[26]: # Evaluate the model
      gru_model.evaluate(X_test, y_test, verbose = 0)
```


```
[26]: 0.10654260218143463
```

# 6. PCA - Rich

Using PCA analysis on our data

# Principal Components Analysis (PCA)

- Goal: Determine the combination of features that explain most of the variation in a data set.
- In our case PCA may help to understand which, if any, Twitter sentiment values contribute to variation in Ethereum price.



Prin. Comp.	Eigenvalue	Variation Percent	Variation Cum Percent
1	6.7127	43.930	43.930
2	2.4661	16.139	60.069
3	1.8951	12.402	72.470
4	1.3600	8.900	81.371
5	1.1768	7.701	89.072
6	0.6942	4.543	93.615
7	0.6601	4.320	97.935
8	0.2237	1.464	99.398
9	0.0881	0.576	99.975
10	0.0028	0.019	99.993
11	0.0008	0.005	99.998
12	0.0003	0.002	100.000



# Principal Components Analysis (PCA)

## Loading Matrix

	Prin1	Prin2	Prin3
Open	0.97643	-0.21404	-0.01543
High	0.97621	-0.21732	-0.00646
Low	0.97813	-0.20903	-0.03619
Close	0.97767	-0.21221	-0.01988
Adj Close	0.97767	-0.21221	-0.01988
Volume	0.62399	0.20969	0.54697
avgGas	-0.23572	-0.39608	-0.00971
nlikes	0.24375	0.41078	0.44950
nretweets	0.34453	-0.09548	-0.09425
compound	0.44309	0.52135	-0.66219
neg	-0.21386	0.15225	0.79418
neu	-0.27968	-0.78327	-0.04183
pos	0.41416	0.73666	-0.39819
Personal Saving %	0.00546	0.63940	0.22501
Close/Last	0.90720	0.02537	0.31987

## Conclusions:

- Features identified in Principal Component 1 are likely to be better predictors of Ethereum closing price.
- Twitter sentiment likely has little predictive value.



# Principal Components Analysis (PCA)

## SAS code (default to old knowledge... time constraints)

```
Principal Components(
    Y(
        :Open, :High, :Low, :Close, :Adj Close, :Volume, :avgGas, :nlikes,
        :nretweets, :compound, :neg, :neu, :pos, :Personal Saving %, :Close/Last"n
    ),
    Estimation Method( "Default" ),
    "on Correlations",
    Eigenvalues( 1 ),
    Loading Matrix( 1 ),
    Factor Analysis( "PC", "SMC", 0, "Varimax" ),
    SendToReport(
        Dispatch( {}, "Loading Matrix", OutlineBox, {Select} ),
        Dispatch(
            {"Factor Analysis: Principal Axis / Varimax"},
            "Prior Communality Estimates:SMC",
            OutlineBox,
            {Close( 0 )}
        ),
        Dispatch(
            {"Factor Analysis: Principal Axis / Varimax"},
            "Eigenvalues of the Reduced Correlation Matrix",
            OutlineBox,
            {Close( 0 )}
        ),
        Dispatch(
            {"Factor Analysis: Principal Axis / Varimax"},
            "Unrotated Factor Loading",
            OutlineBox,
            {Close( 0 )}
        )
    )
);
```

# 7. Summary

Summary of our project

# 8. Postmortem

# Postmortem

## *Difficulties that arose*

Our difficulties and how we dealt with them

- Data collection and cleaning, time consuming
- We ask for help - TAs, Instructor & Reached out to api company on telegram

Additional questions that came up that we didn't answer?

- Understanding model evaluation

What would we research next if we had more time?

- Crypto exploits data and how it affects eth price movements
- Emoji Sentiment analysis

# 9. Questions

Open floor Q&A with the audience





**thank you,**

Special thanks to

- Erik
- Deborah
- Owracle Staff
- US